Campaign Strategy: LatAm Mining Outbound

Neha Girish Kulkarni · FlytBase Outbound BDR Assessment

I built this campaign around one constraint I didn't want to hide from: I have no paid data tools yet, and the whole system has to earn trust one accurate email at a time. Everything below reflects that.

Sequence Design — 3 touches, ~12 days

Day	Channel	Purpose
Day 0	Email	Personalized outreach citing the specific researched signal
Day 4	LinkedIn	Connection request + comment referencing the same signal (not a copy of the email)
Day 9	Email	Short closing note - one line, no re-pitch, just a clear yes/no ask

I kept it to three touches on purpose — these are VP+ industrial executives, not a marketing list, so every touch has to earn its place. LinkedIn sits in the middle deliberately: a connection or comment before the second email lowers the “who is this” resistance without ever feeling automated.

Trigger Logic — what jumps a prospect to the front of the queue

Four signals move an account up the queue: **a recent leadership change** in one of our three goal roles (new hires get a natural 90-day “let me introduce myself” window); **a safety or environmental incident** disclosed in local press — this is the strongest trigger, since it maps directly onto FlytBase's actual value prop of reducing human exposure at hazardous sites; **a public expansion or capex announcement** (more sites means more inspection burden, right when they're thinking about it); and **conference or event attendance** in mining-safety or industrial-automation circles, which signals active budget-seeking rather than theoretical interest.

Personalization at Scale — 10 → 200 accounts

The honest answer is that personalization quality degrades the moment you stop reading every account yourself, so I designed around that instead of pretending it won't happen. Research is cached at the account level, not the contact level, so three contacts at one company share a single research pass instead of three redundant ones. As volume grows, accounts get tiered: the top 20% by ICP fit get a full manual-reviewed pass; the rest draw from a small library of validated signal categories (recent incident, recent expansion, regulatory change), with the company-specific detail still pulled fresh per account — never a swapped name in an unchanged sentence. Human review then scales by sampling a percentage of the tiered tail rather than reviewing everything, which is only safe because an automated QA layer (source-cited claims only, word count, CTA presence) already catches the mechanical misses before a human sees the draft.

Success Metrics — first 2 weeks

Reply rate is the real signal, not open rate — open rate only tells me the subject line worked, not that the message did. Meeting-booked rate is the actual north star. I'm also tracking one metric most people wouldn't: the percentage of contacts landing in “needs manual research” instead of “found.” That's a system-health number, not a campaign number — if it climbs, the problem isn't my messaging, it's that contact discovery is starving, and no amount of email polish fixes that. Broken looks like: normal open rate, near-zero reply rate uniformly across every account — a messaging problem. Working looks like: replies concentrated specifically in the incident- or expansion-triggered accounts, which tells me the trigger logic is driving results, not luck.

What I'd Test First

I'd A/B test whether naming the specific incident in the subject line (e.g. “contractor safety at La Negra”) outperforms a safety-general subject line on the same account with no incident reference. I'm starting with the subject line, not the email body, because differences show up in open-rate data even at the small sample size I'll actually have in week one — body-copy tests need reply-rate volume I won't have yet. If the specific-incident subject wins clearly, it tells me exactly how deep the research has to go before it earns the reviewer's time — which shapes how I build the tiering logic in month two, not just this one campaign.